ERIC HOLT

Garland, TX • 817-500-8000 • eholt723@outlook.com

LinkedIn: https://linkedin.com/in/ericholt723

Portfolio: https://ericholt.dev

GitHub: https://github.com/eholt723

PROFESSIONAL SUMMARY

AI developer specializing in LLM agent orchestration and production Python services, with recent work spanning LangGraph pipelines, Groq inference integration, RAG, and full-stack React/FastAPI deployments. Seven years in technical diagnostics adds a systems-thinking discipline and reliability focus that complements production AI engineering work.

TECHNICAL SKILLS

Programming: Python (Certified), JavaScript, HTML/CSS, C#
AI/ML Frameworks: LangGraph, PyTorch, TensorFlow, Transformers, Scikit-Learn
Data Tools: NumPy, Pandas
Agent & LLM: Groq inference, RAG pipelines, LLM fine-tuning (LoRA), agentic orchestration, reinforcement learning, model evaluation
Web Development: FastAPI, React, Vite, Tailwind CSS, Flask, REST APIs
Tools: Git, Docker, Docker Compose, Render, VS Code, Jupyter, PlatformIO
Embedded & Systems: CAN bus, PLC basics, IoT/Embedded Diagnostics

EDUCATION

DeVry University – Bachelor of Science in Computer Information Systems (Expected 2028)
Specialization: Programming, Machine Learning & Design Technology

CERTIFICATIONS & ACCOMPLISHMENTS

PCEP – Certified Python Programmer, 2025
Dean's List, DeVry University, 2025
Web Scraping with Python, DeVry University, 2025
Vibe Coding with AI, DeVry University, 2025
National Society of Leadership and Success – Inducted Member

TECHNICAL PROJECTS

AgentFlow – HR Document Intelligence (Python, LangGraph, FastAPI, Groq, React, Docker)

- Built a LangGraph multi-step agent pipeline that classifies uploaded HR documents (resumes, cover letters, interview notes, scorecards) and routes each to a specialized LLM analysis node returning structured JSON with skills, risk flags, and hiring recommendations.
- Integrated Groq LLM inference with tenacity retry logic and vision model support for image/handwritten documents; added a deterministic pre-check layer to ensure reliable document classification without LLM dependency.

- Deployed as a single Dockerized service on Render — multi-stage build bundles the React/Vite frontend with the FastAPI backend, with SQLite session persistence and an optional n8n webhook integration point for downstream automation.

AI Tutor - TinyLlama + LoRA (Python, LangGraph, FastAPI, Docker, Azure)

- Built a lightweight AI tutoring backend using a quantized TinyLlama model with a custom LoRA fine-tuned for beginner-level programming instruction.
- Implemented a LangGraph-based orchestration pipeline supporting base vs fine-tuned model comparison and structured tutoring prompts.
- Deployed as a Dockerized FastAPI service on Azure Container Apps with autoscaling-to-zero and CPU-optimized inference via llama.cpp.

Transport LLM-WebGPU (React, Python, RAG, Local Corpus)

- Developed a fully client-side LLM accelerated with WebGPU, enabling on-device inference with zero server dependency.
- Implemented vector search, embeddings, and prompt-augmentation for a custom RAG pipeline.
- Built a React interface controlling weight loading, retrieval visualization, and model inference.

Reinforcement Learning Warehouse Bot (Python, PyTorch, Gymnasium)

- Designed and trained a reinforcement learning agent for warehouse navigation and pick-and-place optimization.
- Implemented custom reward shaping and DQN/PPO training loops.
- Improved simulated route efficiency and reduced idle movement through iterative policy tuning.

Image Classifier Web App (JavaScript, TensorFlow.js)

- Built a browser-based image classifier running entirely on the client.
- Created an interactive UI and deployed for 24/7 availability.
- Demonstrated understanding of ML inference in web environments.

EXPERIENCE

Software Engineer / Full-Stack Developer – Get Set Wedding (Dec 2025–Present)

Get Set Wedding (Startup / Contract)
• Shipping a production web app: auth/onboarding, dashboards, and planning tools.
• Built Supabase persistence (user state + constraints) and hardened login/refresh flows.
• Developing admin features, automations, and subscription-ready infrastructure.

Technical Diagnostics & Systems Specialist – Emergency Vehicle Platforms (2017–2024)

City of Coppell • City of Grapevine • Siddons-Martin Emergency Group